

Dictionary Matching with Fingerprints

Student: Dominic Moylett, Supervisor: Dr. Raphaël Clifford and Dr. Benjamin Sach, Project Type: Research

University of Bristol, Department of Computer Science

Introduction

'Big data' is a common term thrown about nowadays. As computers have become more and more powerful, the amount of information we want to process has also grown in size, and we now need to be smarter when thinking about how we store and process this data.

A common method for reducing space is the streaming model. In this model, we never see the whole input, but instead see it in parts. This is a good choice for pattern matching: the text can be processed one character at a time and thus the space complexity becomes independent of the text.

This project specifically looks at the area of dictionary matching, where you are trying to match one text to many patterns. The Algorithms team have devised a method of solving this problem in $O(\log m)$ time per character and $O(k \log m)$ space, based on Porat and Porat's solution to exact pattern matching in $O(\log m)$ time per character and $O(\log m)$ space. My work is on implementing this algorithm and compare it to other solutions for dictionary matching to see how it performs in practice.

Dictionary Matching Formally

We have a text T of n characters and a set P of k patterns p_1, \dots, p_k with lengths m_1, \dots, m_k . For each index of the text, we output an occurrence at j if $\exists i \in \{1, \dots, k\}$ such that $t_{j-m_i}, \dots, t_j = p_i$.

The typical solution to dictionary matching in the streaming model is the Aho-Corasick algorithm, which solves the problem based on a generalisation of Knuth-Morris-Pratt. It takes $O(1)$ time per character and $O(\sum_{i=1}^k m_i)$ space, and was previously used in `fgrep`.

Current Progress

- ▶ The Aho-Corasick algorithm has been implemented.
- ▶ This new algorithm has been implemented for:
 - ▶ Patterns whose lengths are a power of 2
 - ▶ and patterns which are shorter than k .
- ▶ Current work is focusing on long patterns which are repetitive.

Pattern Matching in Less Space than the Pattern?

Sublinear space can be achieved by using a fingerprint function developed by Karp and Rabin for a string of characters t_1, \dots, t_n , a prime number p and a randomly selected r :

$$\Phi_{p,r}(t_1, \dots, t_n) = \sum_{i=1}^n r^i t_i (\text{mod } p)$$

These fingerprints compress the amount of space required to store text, and can be modified to match changes in the underlying strings.

Future Work

- ▶ Finish the work on long patterns which are repetitive.
- ▶ Implement the algorithm for long patterns.
- ▶ Compare the performance of the new algorithm to Aho-Corasick.
- ▶ Implement optimisations to the algorithm.

